

FlowMind

Turn Text Chaos Into Execution Clarity

Built with Chutes AI

Chutes AI Hackathon 2025

The Problem



Notes pile up, unread

Meeting transcripts and project briefs sit in text form — no structure, no action items, no owners.



Tasks get buried in walls of text

Critical decisions and blockers are hidden inside paragraphs of prose, easy to miss.



Manual extraction wastes hours

Teams spend valuable time reading, highlighting, and copying tasks into project management tools.

Meet FlowMind

Paste any unstructured text — meeting notes, project briefs, random notes — and FlowMind's 4-layer AI pipeline transforms it into a structured work card and a drag-and-drop Kanban board, ready for immediate execution.

1

Raw Text

Meeting notes
Project briefs
Any unstructured text

2

AI Pipeline

4 intelligent
layers

3

Work Card

Tasks · Timeline
Risks · Next Action

4

Kanban Board

Drag-drop task
management

4-Layer AI Pipeline

Powered by Qwen3.6-27B-TEE on Chutes AI

1

Understanding

Classifies input type (Meeting / Project Brief / Notes) with a confidence score. Extracts participants, topics, and key context.

2

Processing

Pulls out structured data: task items with owners & deadlines, risks, milestones, and key decisions.

3

Synthesis

Creates an executive summary and selects the single highest-priority next action for the team.

4

Formatting

Validates the schema, fills missing fields with smart fallbacks, and returns a clean WorkCard.

Work Card Output

One input. One complete work card. Everything your team needs to execute.



Summary

AI-generated executive summary of the entire input



Tasks

Action items with assigned owner and deadline



Timeline

Milestones and key dates extracted automatically



Risks

Blockers and concerns the AI identified



Next Action

The single most important thing to do right now

AI Work Card — Sample

Summary

Q3 planning review: launch pushed to Oct 15. 3 blockers identified. Marketing & Eng aligned on scope.

Tasks (3)

Fix API latency — Shan, Oct 8

Update design mocks — Jay, Oct 6

Review compliance doc — Mei, Oct 10

Next Action

Fix API latency — highest urgency

Risks (2)

Auth service latency spike

Design review bottleneck

Drag-and-Drop Kanban Board

Tasks from the work card auto-populate the board. Edit, move, track — no manual setup.

To Do

Fix API latency

Shan · Oct 8 · Urgent



In Progress

Update design mocks


Jay · Oct 6 · Mild



Done ✓


Kick-off meeting notes

Team · Oct 1 · Done



Review compliance doc

Mei · Oct 10 · Mild



Transparency Built In

Watch each AI layer activate in real time. No black box — full confidence visibility.

Layer 1 · Understanding

Done



Layer 2 · Processing

Done



Layer 3 · Synthesis

Running...



Layer 4 · Formatting

Pending

Input classified as: Meeting Transcript · Confidence: 92%

Built With

Frontend

- React 18 + Vite 5
- TailwindCSS 4
- Drag-and-Drop Kanban
- Real-Time Pipeline Visualizer
- OAuth PKCE Auth Flow

Backend + AI

- FastAPI + Uvicorn
- Pydantic v2 schemas
- Chutes AI (Qwen3.6-27B-TEE)
- 4-Layer Orchestration pipeline
- OpenAI-compatible fallback

Why Chutes AI?



Privacy by Design (TEE)

Qwen3.6-27B runs in a Trusted Execution Environment — users' meeting notes and project briefs stay private. Critical for enterprise data.



User-Billable OAuth

Each user authenticates with their own Chutes account. AI costs bill directly to them — no backend key required. Perfect for multi-user apps.



High-Quality Reasoning

27B parameters with strong instruction-following made our 4-layer structured extraction reliable and accurate from day one.

FlowMind

From chaos to clarity — in seconds.



GitHub Repo

[github.com/\[your-username\]/flowmind](https://github.com/[your-username]/flowmind)



Demo Video

[youtube.com/\[paste your unlisted link\]](https://youtube.com/[paste your unlisted link])